

Native-Task Performance Test Report

Intel Software

Wang, Huafeng, Huafeng.wang@intel.com

Zhong, Xiang, xiang.zhong@intel.com

1. Background

2. Related Work

3. Preliminary Experiments

3.1 Experimental Environment

Workbench	Peculiarity
Wordcount	CPU-intensive
Sort	IO-intensive
DFSIO	IO-intensive
Pagerank	Map :CPU-intensive Reduce :IO-intensive
Hivebench-Aggregation	Map :CPU-intensive Reduce :IO-intensive
Hivebench-Join	CPU-intensive
Terasort	Map :CPU-intensive Reduce : IO-intensive
K-Means	Iteration stage: CPU-intensive Classification stage: IO-intensive
Nutchindexing	CPU-intensive & IO-intensive

Cluster settings

Hadoop version	1.0.3-Intel (patched with native task)
Cluster size	4
Disk per machine	7 SATA Disk per node
Network	GbE network
CPU	E5-2680(32 core per node)
L3 Cache size	20480 KB
Memory	64GB per node
Map Slots	$3*32+1*26=122$
Reduce Slots	$3*16+1*13=61$


Job Configuration

io.sort.mb	1GB
compression	Enabled
Compression algo	snappy
Dfs.block.size	256MB
io.sort.record.percent	0.2
Dfs replica	3

3.2 Performance Metrics

	Data before compression	size	Data after compression	size	Native job run time(s)	Original job run time(s)	Job performance improvement	Map stage performance
--	-------------------------------	------	------------------------------	------	---------------------------------	-----------------------------------	-----------------------------------	-----------------------------

)		ment	improve ment
Wordcount	1TB	500GB	1523.43	3957.11	159.8%	159.8%
Sort	500GB	249GB	2662.43	3066.97	15.2%	45.4%
DFSIO-Read	1TB	NA	1249.68	1384.52	10.8%	26%
DFSIO-Write	1TB	NA	6639.22	7165.97	7.9%	7.9%
Pagerank	Pages:500M Total:481GB	217GB	6105.71	11644.63	90.7%	133.8%
Hive-Aggregation	Uservisits:5 G Pages:600M Total:820GB	345GB	1113.82	1662.74	49.3%	76.2%
Hive-Join	Uservisits:5 G Pages:600M Total:860GB	382GB	1107.55	1467.08	32.5%	42.8%
Terasort	1TB	NA	4203.35	6360.49	51.3%	109.1%
K-Means	Clusters:5 Samples:2G Inputfilesec ondample:4 00M Total:378GB	350GB	5734.11	7045.62	22.9%	22.9%
Nutchindexing	Pages:40M 22G	NA	4387.6	4600.56	4.9%	13.2%





3.3 Results

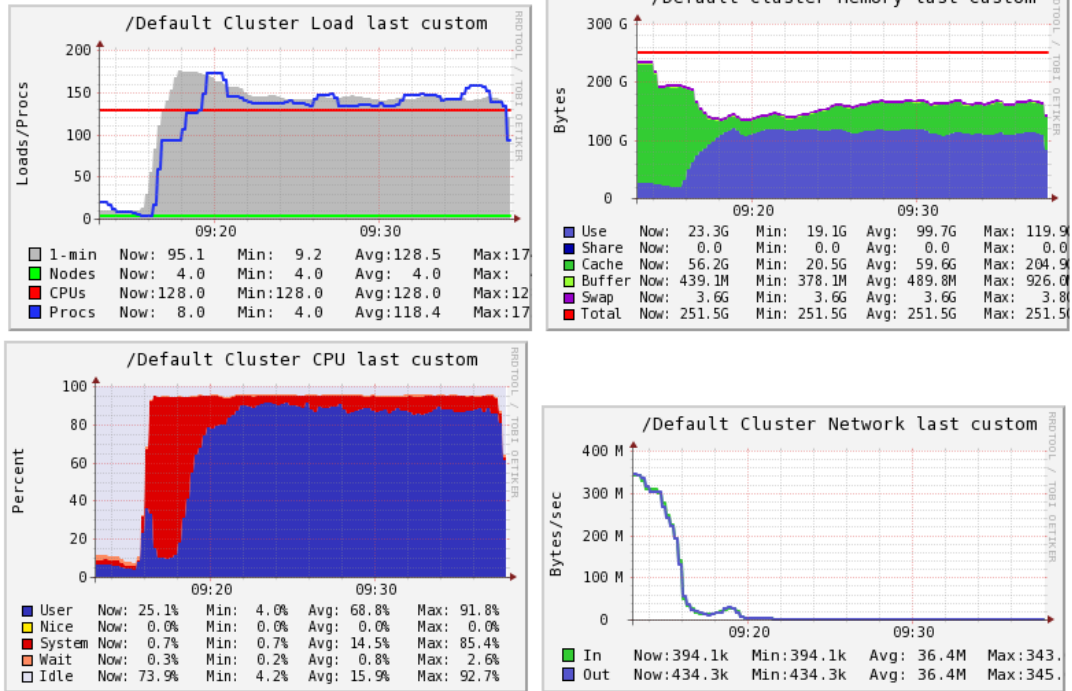
3.3.1 Wordcount

➤ Job Details:

Name	Maps	Reducers
wordcount	1984	22



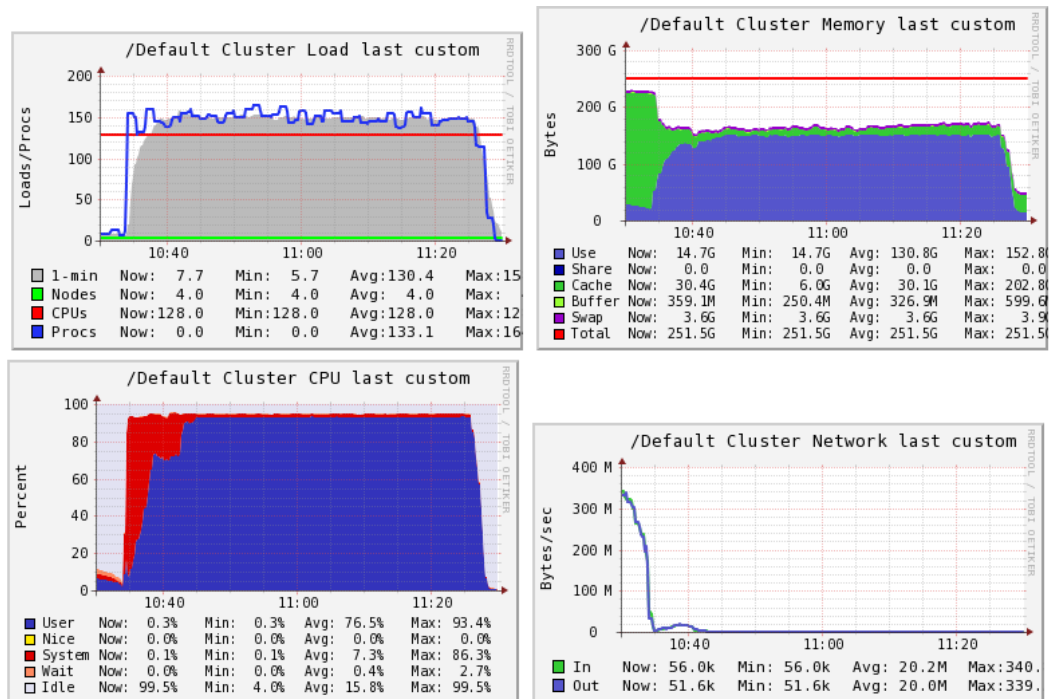
➤ Native-Task running state:



Start time: 9:14

Finish time: 9:37

➤ Original running state:



Start time: 10:32

Finish time: 11:28



Analysis

Wordcount is a CPU-intensive workload and it's map stage run through the whole job. So the native-task has a huge performance improvement.

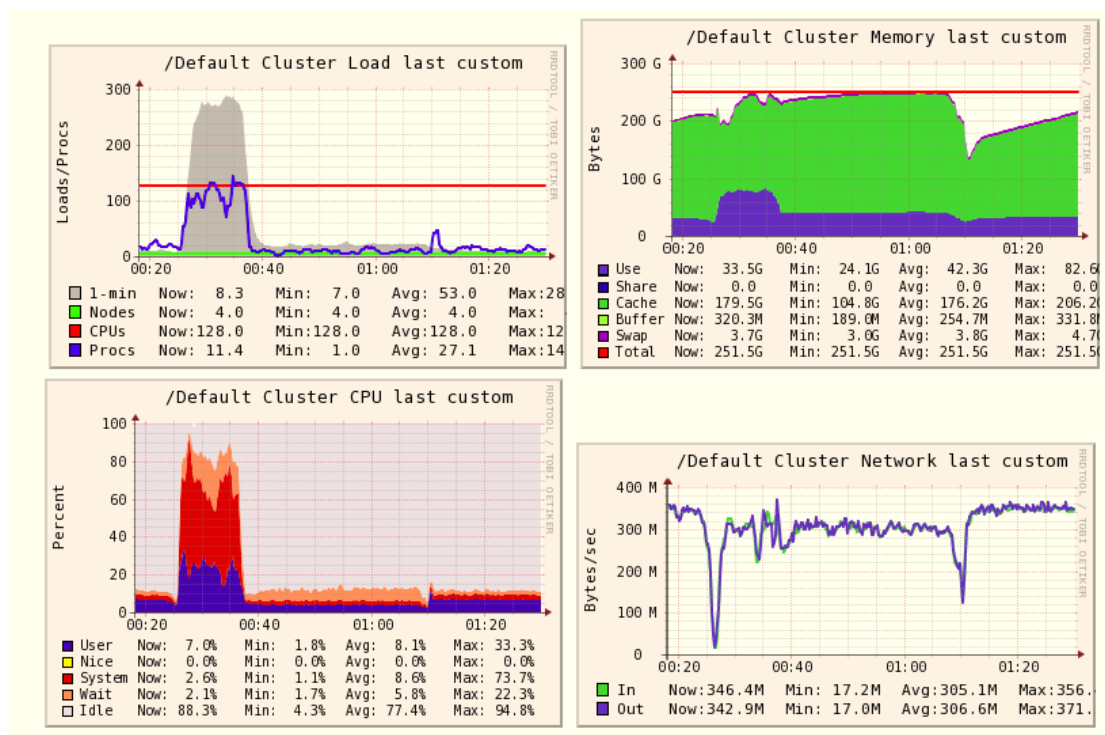
3.3.2 Sort

➤ Job Details:

Name	Maps	Reducers
sorter	1024	33



➤ Native-Task running state:



Start time :00:25

Finish time :1:10



Analysis

Sort is IO-intensive at both map and reduce stage. We can see that it's reduce time occupy the most of whole job running time, because of that, the performance improvement is limited.

3.3.3 DFSIO-Read

➤ Job Details:

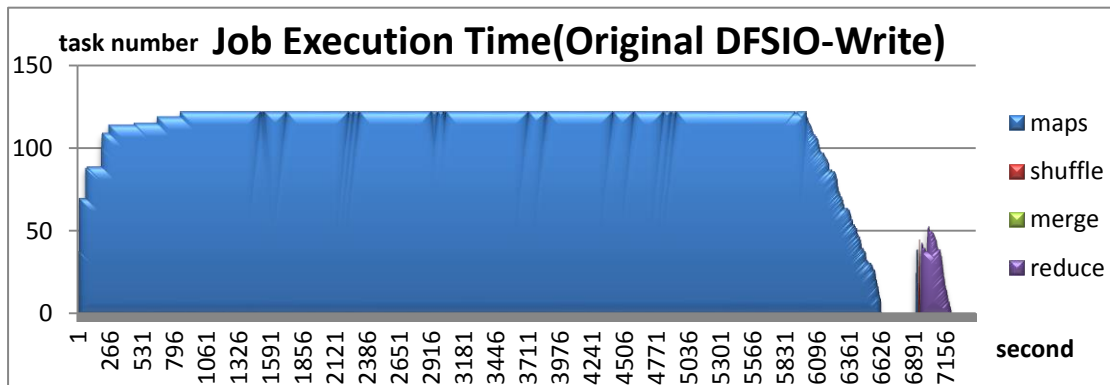
Name	Maps	Reducers
Datatools.jar	256	1
Result Analyzer	50	63




3.3.4 DFSIO-Write

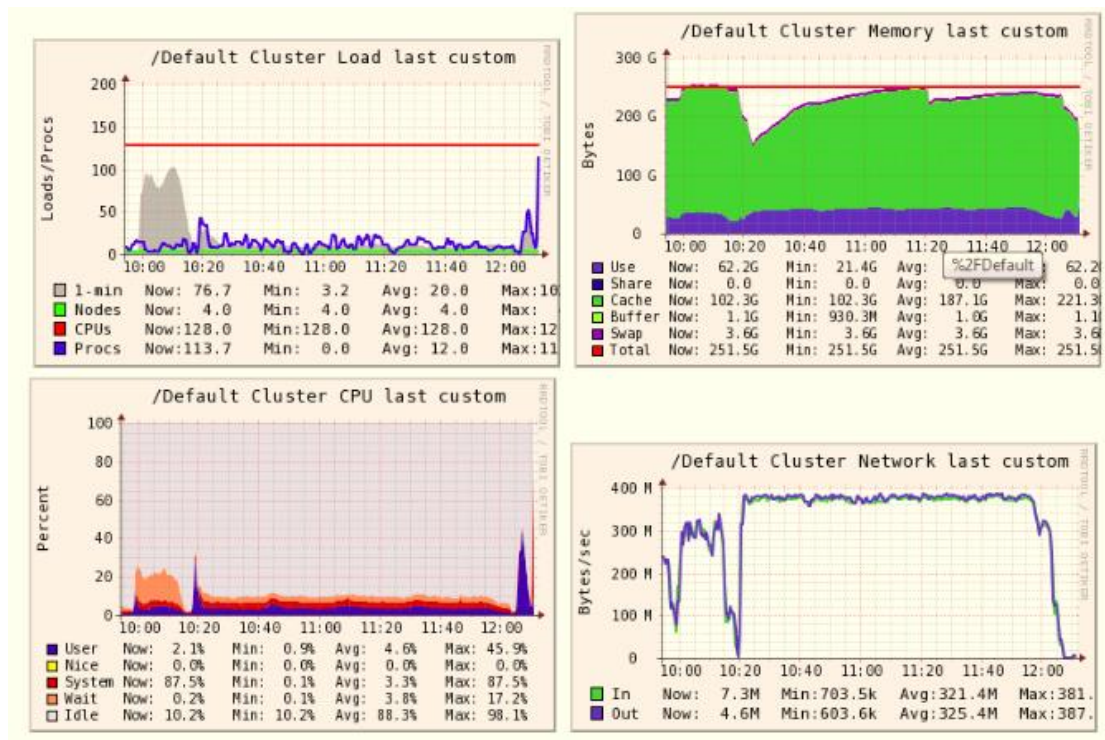
➤ Job Details:

Name	Maps	Reducers
Datatools.jar	512	1
Result Analyzer	50	63





➤ Native-Task running state:




Aggregation start time: 9:58

Aggregation finish time: 10:19

Join start time: 10:19

Join finish time: 12:10

➤ Original running state:



Aggregation start time: 20:22

Aggregation finish time: 20:46

Join start time: 20:46

Join finish time: 22:45



Analysis

DFSIO is IO-intensive both at read and write stage. It's bottleneck is network bandwidth so the performance improvement is limited.

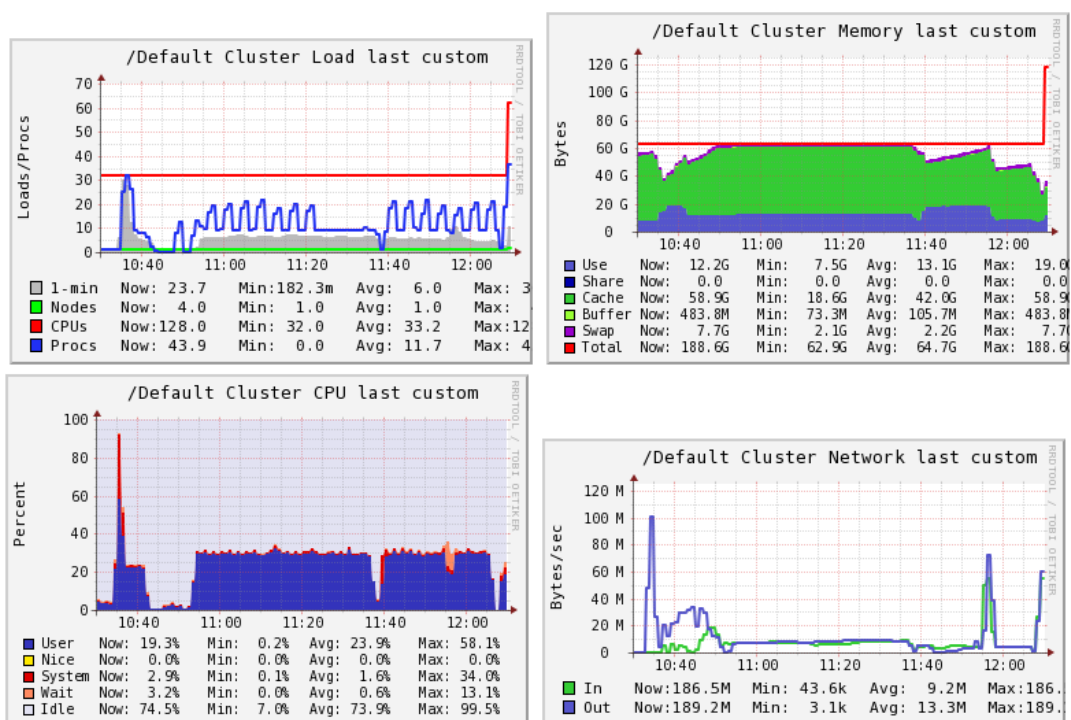
3.3.5 Pagerank

➤ Job Details:

Name	Maps	Reducers
Pagerank_Stage1	163	33



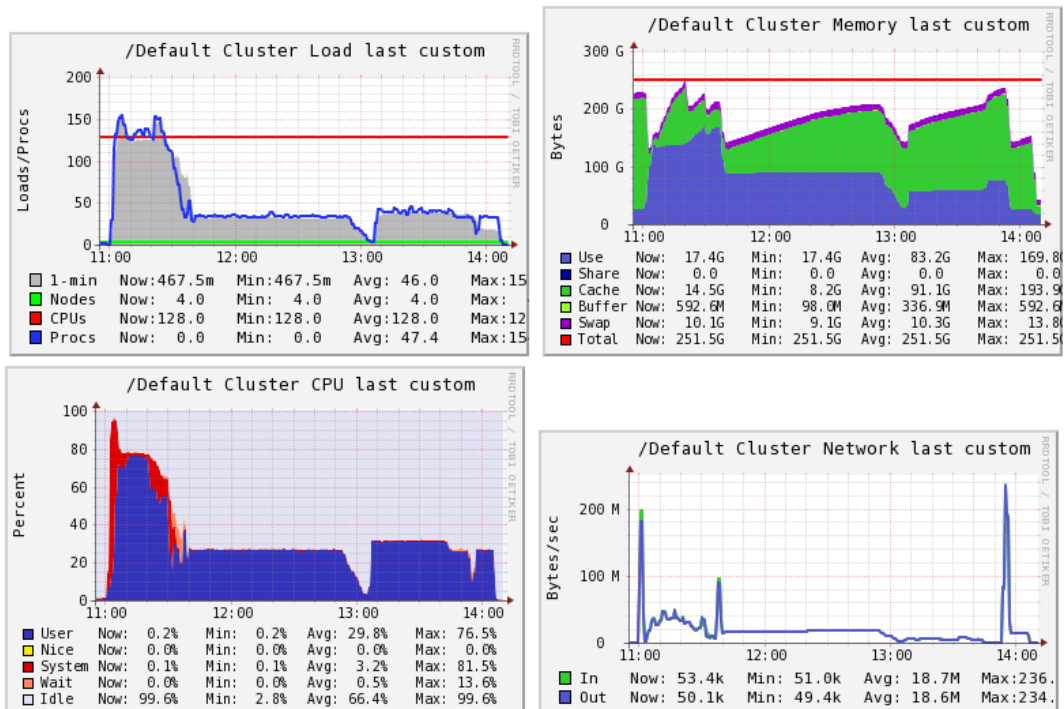
➤ Native-Task running state:



Start time: 10:33

Finish time: 12:08

➤ Original running state:



Start time: 10:59

Finish time: 14:06

Analysis



Pagerank is a CPU-intensive workload and its map stage take about 50% of the whole job running time. So the performance improvement is obvious.

3.3.6 Hive-Aggregation

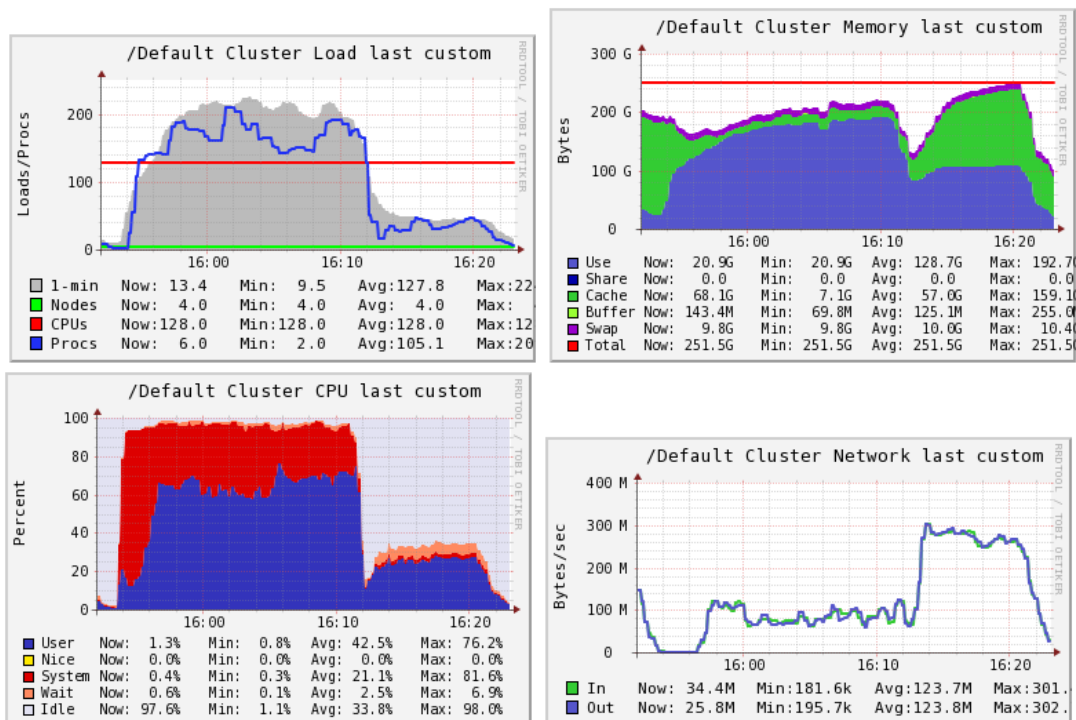
➤ Job Details:

Name	Maps	Reducers
INSERT OVERWRITE TABLE	1386	33

uservisits...sourceIP(Stage-1)



➤ Original running state:



Start time: 15:52

Finish time :16:22



Analysis

Hive-Aggregation is CPU-intensive at map stage and IO-intensive at reduce stage. It's map stage occupy the most of running time and when it comes to reduce stage, network bandwidth limits the performance. So the performance improvement at map stage is obvious.





3.3.7 Hive-Join

➤ Job Details:

Name	Maps	Reducers
INSERT OVERWRITE TABLE rankings_uservisi...1(Stage-1)	1551	33
INSERT OVERWRITE TABLE rankings_uservisi...1(Stage-2)	99	33
INSERT OVERWRITE TABLE rankings_uservisi...1(Stage-3)	99	1
INSERT OVERWRITE TABLE rankings_uservisi...1(Stage-4)	1	1



➤ Original running state:



Start time: 16:32

Finish time :16:58



Analysis

Hive-join is a CPU-intensive workload and its map stage takes a high percent of whole running time. So we can see at map stage, the performance is improved by native-task.

3.3.8 Terasort

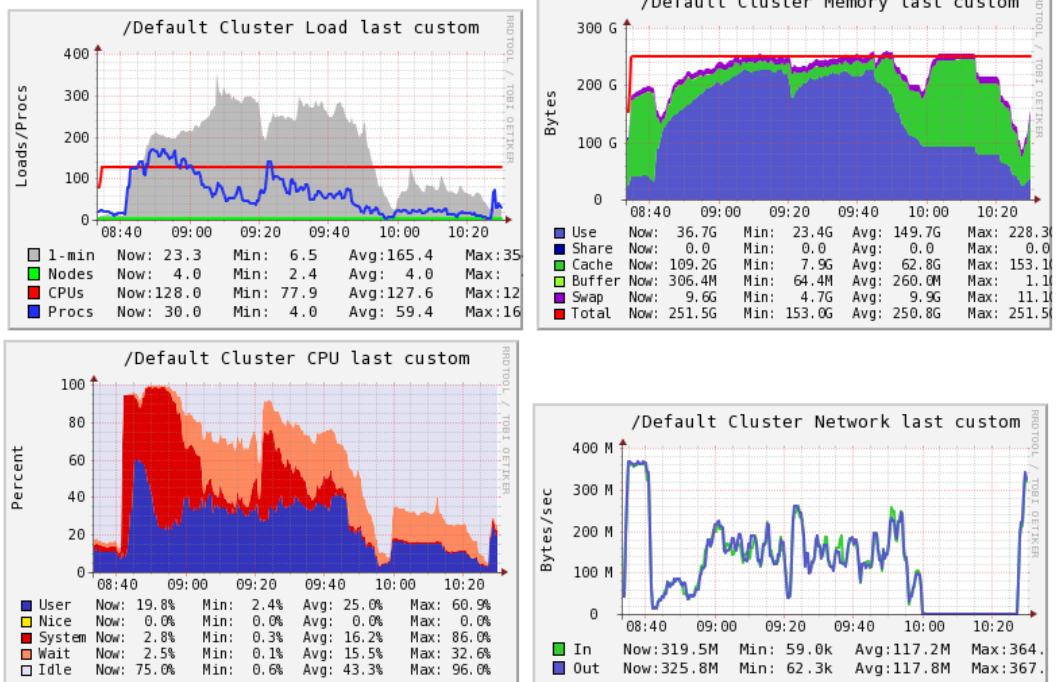
➤ Job Details:

Name	Maps	Reducers
Terasort	458	34



➤ Native-Task running state:

➤ Original running state:



Start time: 8:39

Finish time: 10:24

Analysis

Terasort is CPU-intensive at map stage and IO-intensive at reduce stage. Its map stage occupies the majority of the running time so there is a huge performance improvement at map stage.

3.3.9 K-Means

➤ Job Details:

Name	Maps	Reducers
Cluster Iterator running iteration 1 ...	1400	63
Cluster Iterator running	1400	63

iteration 2 ...

Cluster Iterator running	1400	63
--------------------------	------	----

iteration 3 ...

Cluster Iterator running	1400	63
--------------------------	------	----


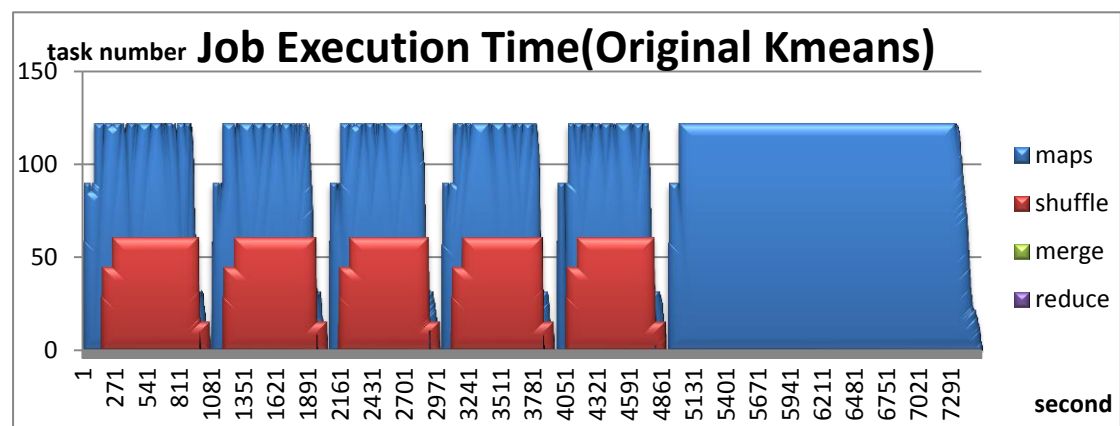
iteration 4 ...

Cluster Iterator running	1400	63
--------------------------	------	----

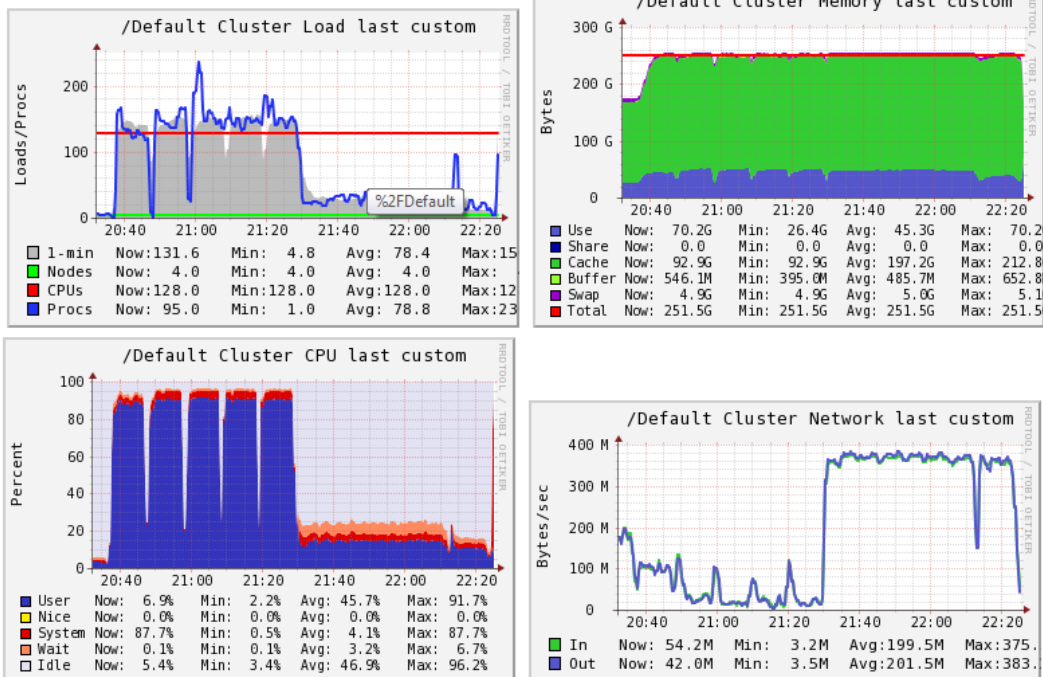
iteration 5 ...

Cluster	Classification	1400	0
---------	----------------	------	---

Driver running



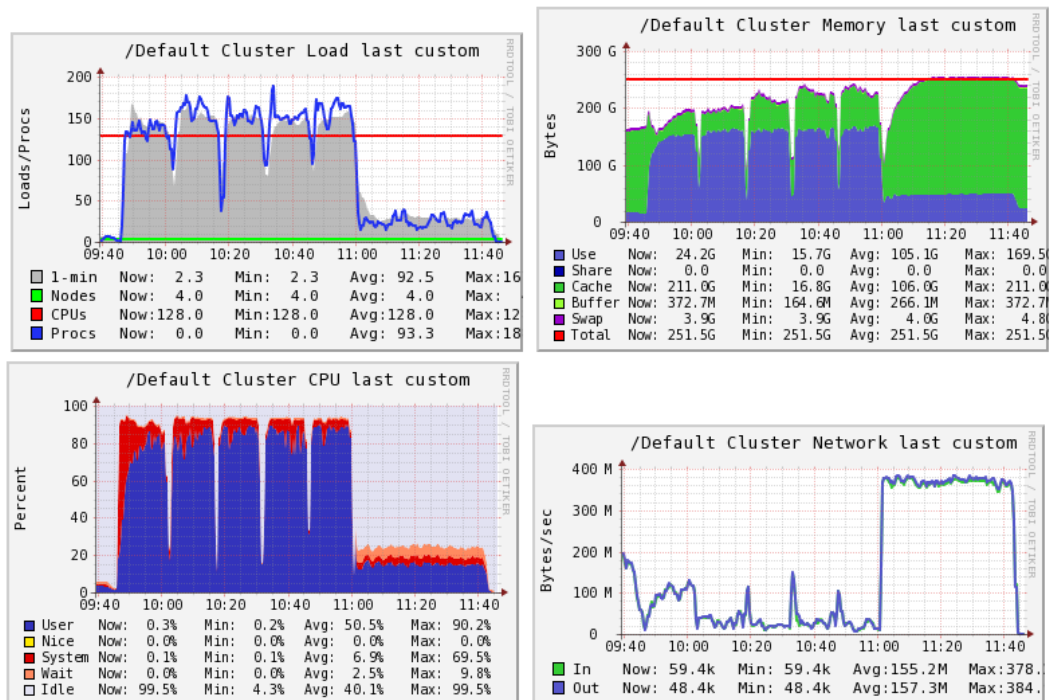
➤ Native-Task running state:



Start time: 20:38

Finish time: 22:23

➤ Original running state:



Start time: 9:43

Finish time: 11:41



Analisis

From the running state graph, we can see that the former 5 iteration is CPU-intensive and the last classification stage is IO-intensive. The two stages almost equally split the whole running time. So the performance improvement at map stage is evident.





3.3.10 Nutchindexing

➤ Job Details:

Name	Maps	Reducers
index-lucene	1171	133
/HiBench/Nutch/Input/indexes		







➤ Native-Task running state:



Start time: 17:26

Finish time: 18:40

➤ Original running state:



Start time: 18:40

Finish time: 19:56


Analysis

Nutchindexing is CPU-intensive at map stage but the reduce stage take the majority of whole running time. So the performance improvement is not so huge.


3.4 Other related results

3.4.1 Cache miss hurts Sorting performance

- Sorting time **increase rapidly** as cache miss rate increase
- We divide a large buffer into several memory unit.
- BlockSize is the size of memory unit we doing the sorting.




3.4.2 Compare with BlockMapOutputBuffer



- **70% faster** than BlockMapOutputBuffer collector.
- BlockMapOutputBuffer **supports ONLY BytesWritable**

3.4.3 Effect of JVM reuse


- 4.5% improvement for Original Hadoop, 8% improvement for Native-Task



4 nodes, 4 map slots per node.


3.4.4 Hadoop don't scale well when slots number increase

- 4 nodes(32 core per node), 16 map slots max, **CPU, memory, disk are NOT fully used.**
- **Performance drops unexpectedly when slots# increase**



3.4.5 Native-Task mode: full task optimization

- 2x faster further for Native-Task full time optimization, compared with native collector



4. Conclusions